

Assignment 3: EN4553 (Machine Vision)

Thalagala B.P. 180631J

Submitted on February 19, 2023

1 (a)

(5 points) Complete the following Python function to read the dataset as a set of arrays. (Hint: you may use the `np.loadtxt()` function.)

```
[3]: def load_dataset( src_dir: str ) -> Tuple[np.ndarray, np.ndarray, np.
    ↳ ndarray, np.ndarray, np.ndarray]:
    """Load the dataset as a set of numpy arrays.
    Args:
        src_dir: Directory where dataset files are stored.
    Returns:
        (x_train, y_train, x_val, y_val, x_test) tuple where each array is_
    ↳ one dimensional.
    """
    # reading data from the files
    x_train = np.loadtxt(src_dir + "/x_train.txt")
    y_train = np.loadtxt(src_dir + "/y_train.txt")
    x_val   = np.loadtxt(src_dir + "/x_val.txt")
    y_val   = np.loadtxt(src_dir + "/y_val.txt")
    x_test  = np.loadtxt(src_dir + "/x_test.txt")

    # returning as a tuple
    return x_train, y_train, x_val, y_val, x_test
```

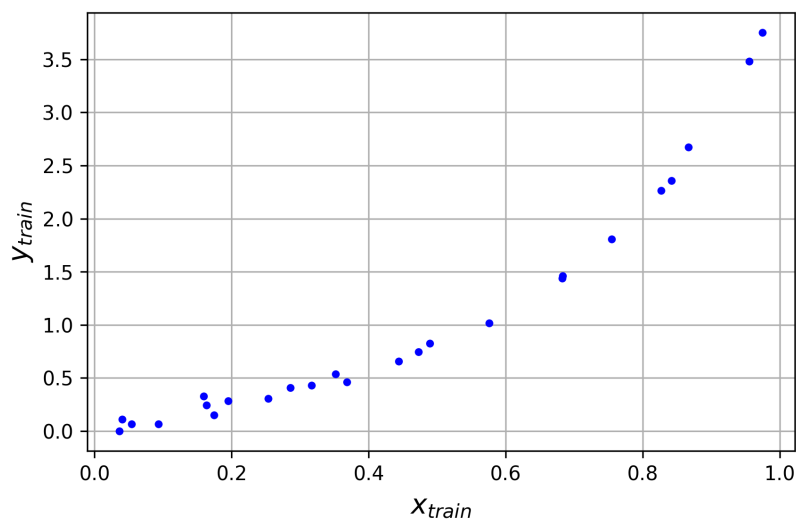


Figure 1: Training Data

Training Dataset size: (25,)

2 (b)

i. (5 points) Implement the following function to make input features for the above linear regression model.

```
[5]: def get_features(x: np.ndarray, n: int) -> np.ndarray:
    """Creates n-th degree polynomial features for the given vector x.
    Example usage:
    get_features(np.array([1.0, 2.0, 3.0]), 3) outputs
    np.array([ [ 1., 1., 1.],
               [ 2., 4., 8.],
               [ 3., 9., 27.]])

    Args:
        x: A numpy array of shape (num_examples, ) or (num_examples, 1).
        n: The degree of the polynomial features.

    Returns:
        A matrix of shape (num_examples, n) where the j-th column is equal to
        the vector x raised, elementwise, to the power j.
    """

    # create an array of powers
    powers = np.arange(1, n+1)

    # reshape x to have the shape (num_examples, 1) to be compatible with np.
    ↪power broadcasting
    x = x.reshape(-1, 1)

    # construct the design matrix of shape (num_examples, n)
    design_matrix = np.power(x, powers)

    # return the design matrix
    return design_matrix
```

ii. (10 points) Use the above function to complete the following implementation. (Hint: you may use the `sklearn.linear_model.LinearRegression` class.)

```
[6]: def fit_and_evaluate(
    x_train: np.ndarray, y_train: np.ndarray,
    x_val: np.ndarray, y_val: np.ndarray,
    n: int
) -> Tuple[float, float]:

    """Fits an n-th degree polynomial and outputs train and validation MSE.
    Fits a linear regression model  $y = \sum_{i=1}^n w_i x^i$  to the given train
    set and outputs the mean-squared-error (MSE) on train and validation_
    ↪sets.

    Args:
    x_train: Input features for the train set. Has shape (num_train, )
    y_train: Targets (labels) for the train set. Has shape (num_train, )
    x_val: Input features for the validation set. Has shape (num_val, )
    y_val: Targets (labels) for the validation set. Has shape (num_val, )
    n: The degree of the polynomial fit. See the above equation.

    Returns:
    (train_mse, val_mse), tuple of MSE on train and validation sets.
    """

    # Generating polynomial features using the defined function
    x_poly_train = get_features(x_train, n)
    x_poly_val   = get_features(x_val, n)

    # Fitting a linear regression model to the training set
    lin_reg = LinearRegression()
    lin_reg.fit(x_poly_train, y_train)

    # Generate model predictions for the train set and calculate the MSE.
    y_predict_train = lin_reg.predict(x_poly_train)
    train_mse = mean_squared_error(y_train, y_predict_train)

    # Similarly, calculate the MSE on the val set.
    y_predict_val = lin_reg.predict(x_poly_val)
    val_mse = mean_squared_error(y_val, y_predict_val)

    return train_mse, val_mse
```

3 (c)

(5 points) Use the above function to calculate and plot train and validation MSEs against $n = 1, 2, \dots, 10$. Include this graph in your answer sheet. Which n value would you pick for your final model?

```
[7]: highest_order = 10
degrees = np.arange(1, highest_order+1)
Train_MSE = np.zeros((highest_order,1))
Val_MSE = np.zeros((highest_order,1))

for i in range(1, highest_order +1):

    train_mse, val_mse = fit_and_evaluate(
        dataset[0], dataset[1], #  $x_{train}$ ,  $y_{train}$ 
        dataset[2], dataset[3], #  $x_{val}$ ,  $y_{val}$ 
        i) # The degree of the polynomial to fit

    Train_MSE[i-1] = train_mse
    Val_MSE[i-1] = val_mse

plt.figure(figsize=(6, 4))
plt.plot(degrees, Train_MSE, "r--o", label="Train MSE")
plt.plot(degrees, Val_MSE, "b--o", label="Validation MSE")
plt.legend(loc="upper right")
plt.xlabel("$Degree$"); plt.xticks(degrees)
plt.ylabel("$MSE$", rotation=0); plt.yscale("log")
plt.grid()
save_fig("train_val_mse")
plt.show()
```

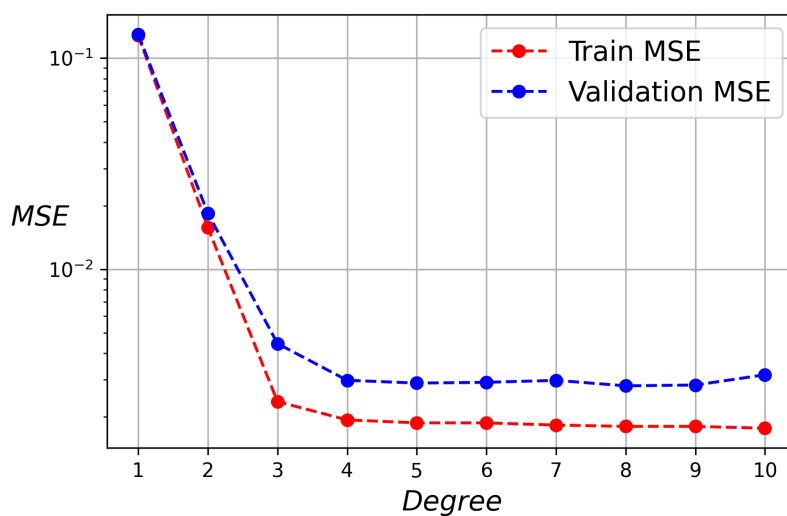


Figure 2: Train MSE and Validation MSE

3.1 (d)

(5 points) Use the model selected above to make predictions on the test set. Include your predictions in a file named <your-index-number>_y_predict_test.txt. For example, 180000X_y_predict_test.txt.

Answer : According to the learning curve, 4th order polynomial fits the data well.

```
[8]: selected_order = 4

# Generating polynomial features using the user defined function
x_poly_train = get_features(dataset[0], selected_order)

# Fitting a linear regression model to the training set
lin_reg = LinearRegression()
lin_reg.fit(x_poly_train, dataset[1])

# predicting on the test data set and save in a file
x_poly_test = get_features(dataset[4], selected_order)
y_predict_test = lin_reg.predict(x_poly_test)
np.savetxt("18063J_y_predict_test.txt", y_predict_test)
```

4 Visualization

```
[9]: # training data
plt.figure(figsize=(6, 4))
plt.plot(dataset[0], dataset[1], "g*", label="Training Data")

# data to plot the polynomial
data_points = 100
x = np.linspace(np.min(dataset[0]), np.max(dataset[0]), data_points).
    ↪ reshape(data_points, 1)
x_poly = get_features(x, selected_order)
y = lin_reg.predict(x_poly)
plt.plot(x, y, "r-", linewidth=2, label="Predictions")

plt.legend(loc="upper left")
plt.xlabel("$X$")
plt.ylabel("$Y$", rotation=0)
plt.grid()
save_fig("predictions")
plt.show()
```

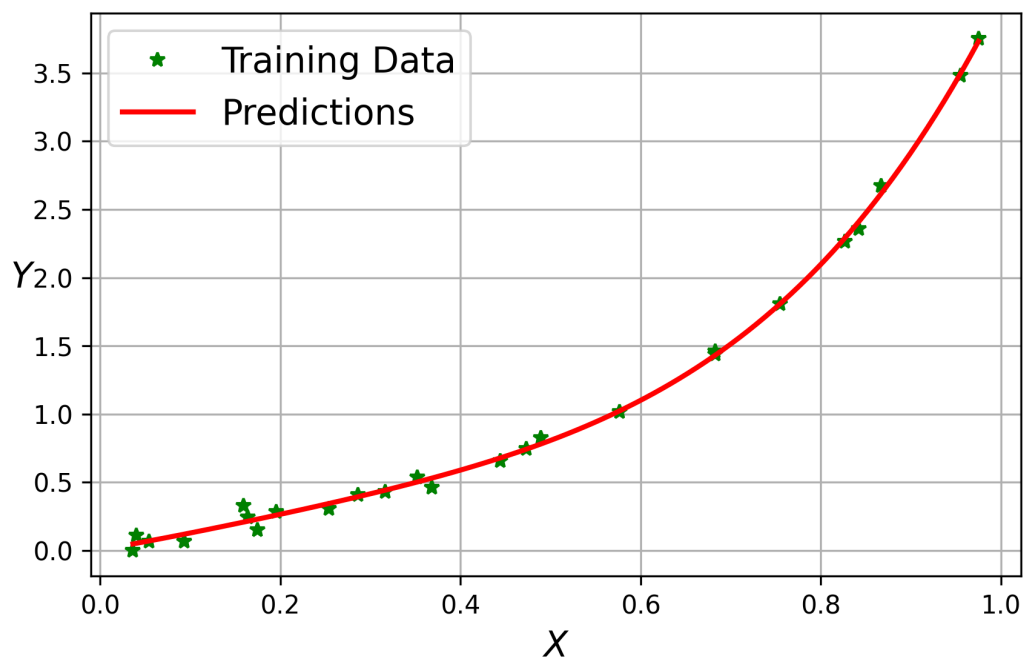


Figure 3: Predictions and the Training Data